

Audio Classification using CNN

In this notebook, we will exploit the features obtainable with torch audio to build a simple audio classifier.

Let's start by importing required libraries.

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchaudio
import sys

import matplotlib.pyplot as plt
import IPython.display as ipd

from tqdm import tqdm
```

It's better to use GPU computing to speed-up training.

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)

cuda
```

Dataset

For this experiment, we are going to use the "SPEECHCOMMANDS" dataset, available on the basic torchaudio datasets (other datasets available here: <https://pytorch.org/audio/stable/datasets.html>).

Let's create a dedicated class to represents a portion of the SC dataset (either training, validation, or testing).

```
from torchaudio.datasets import SPEECHCOMMANDS
import os

class SubsetSC(SPEECHCOMMANDS):
    def __init__(self, subset: str = None):
        super().__init__(("./", download=True))

    def load_list(filename):
        filepath = os.path.join(self._path, filename)
        with open(filepath) as fileobj:
            return [os.path.normpath(os.path.join(self._path,
```

```

line.strip())) for line in fileobj]

    if subset == "validation":
        self._walker = load_list("validation_list.txt")
    elif subset == "testing":
        self._walker = load_list("testing_list.txt")
    elif subset == "training":
        excludes = load_list("validation_list.txt") +
load_list("testing_list.txt")
        excludes = set(excludes)
        self._walker = [w for w in self._walker if w not in
excludes]

# Create training and testing split of the data. We do not use
validation in this tutorial.
train_set = SubsetSC("training")
test_set = SubsetSC("testing")

```

100%|██████████| 2.26G/2.26G [01:58<00:00, 20.4MB/s]

Add support functions (e.g., function to play audio).

```

from IPython.display import Audio

def plot_waveform(waveform):
    waveform_numpy = waveform.t().numpy()
    plt.plot(waveform_numpy);
    plt.show(block=False)

def play_audio(waveform, sample_rate):
    waveform = waveform.numpy()

    num_channels, num_frames = waveform.shape
    if num_channels == 1:
        display(Audio(waveform[0], rate=sample_rate)) # for stereo audio
(two channels)
    elif num_channels == 2:
        display(Audio((waveform[0], waveform[1])), rate=sample_rate))
    else:
        raise ValueError("Waveform with more than 2 channels are not
supported.")

```

Let's explore a tuple of the dataset. It contains the audio waveform and sample_rate, the target label (supervised), the speaker identifier, and the id of the utterance (command).

```

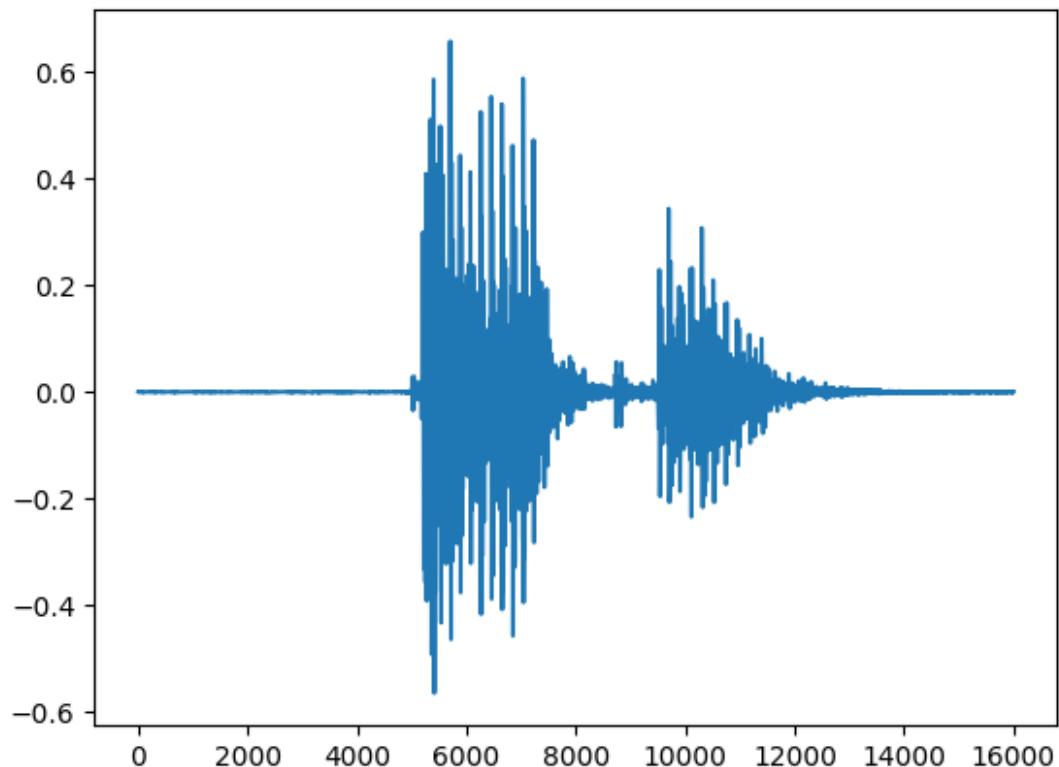
waveform, sample_rate, label, speaker_id, utterance_number =
train_set[2]

```

```
print(waveform, waveform.size())
print("Sample rate: ", sample_rate)

plot_waveform(waveform)

tensor([[-6.1035e-05, -1.8311e-04, -2.4414e-04, ..., -9.1553e-05,
        2.4414e-04, 4.2725e-04]]) torch.Size([1, 16000])
Sample rate: 16000
```



```
labels = sorted(list(set(datapoint[2] for datapoint in train_set)))
labels

['backward',
 'bed',
 'bird',
 'cat',
 'dog',
 'down',
 'eight',
 'five',
 'follow',
 'forward',
 'four',
 'go',
 'happy',
```

```

'house',
'learn',
'left',
'marvin',
'nine',
'no',
'off',
'on',
'one',
'right',
'seven',
'sheila',
'six',
'stop',
'three',
'tree',
'two',
'up',
'vesual',
'wow',
'yes',
'zero']

print('Total number of entries in training set :', (len(train_set)))
print('Total number of entries in test set :', (len(test_set)))

Total number of entries in training set : 84843
Total number of entries in test set : 11005

```

We define a resample transformation, we are going to use it later in the training phase. With this we can apply re-sampling using *transform*.

```

import torchaudio.transforms as T
new_sample_rate = 8000
transform = T.Resample(sample_rate, new_sample_rate)

#let's test it on the previously extracted waveform (see previous
#cell)
transformed = transform(waveform)
play_audio(transformed,new_sample_rate)

<IPython.lib.display.Audio object>

```

Let's now build two simple support functions. The former is to obtain the position of a command (word) in the list of labels - note it will return a tensor (see later use). The latter simply returns the command (label) given an input index.

```

def label_to_index(word):
    # Return a tensor with the position of the word into the list of
    # labels.
    return torch.tensor(labels.index(word))

def index_to_label(index):
    # Return the word in the sorted list of labels
    return labels[index]

word = "visual"
index = label_to_index(word)
word_retrieved = index_to_label(index)

print(word, "-->", index, "-->", word_retrieved)
visual --> tensor(31) --> visual

```

Building DataLoader

We need to prepare batches of audio tensors (i.e., waveforms). To do so, we create a collate_fn function (see <https://pytorch.org/docs/stable/data.html#working-with-collate-fn>). Also, audio samples have different lenght (duration) and, therefore, we need to make all of them aligned. So we use the pad_sequence function and add zero padding to the tensors to make them aligned. We need to engineer a little bit the output tensor in light of the output produced by the rnn.pad_sequence function. Input to our function is a list of waveforms.

```

def pad_sequence(batch):
    # Make all tensor in a batch the same length by padding with zeros
    batch = [item.t() for item in batch]
    batch = torch.nn.utils.rnn.pad_sequence(batch, batch_first=True,
padding_value=0.)
    return batch.permute(0, 2, 1)

def collate_fn(batch):
    # A data tuple has the form:
    # waveform, sample_rate, label, speaker_id, utterance_number

    tensors, targets = [], []

    # Gather in lists, and encode labels as indices
    for waveform, _, label, *_ in batch:
        tensors += [waveform]
        targets += [label_to_index(label)]

    # Group the list of tensors into a batched tensor
    tensors = pad_sequence(tensors)

```

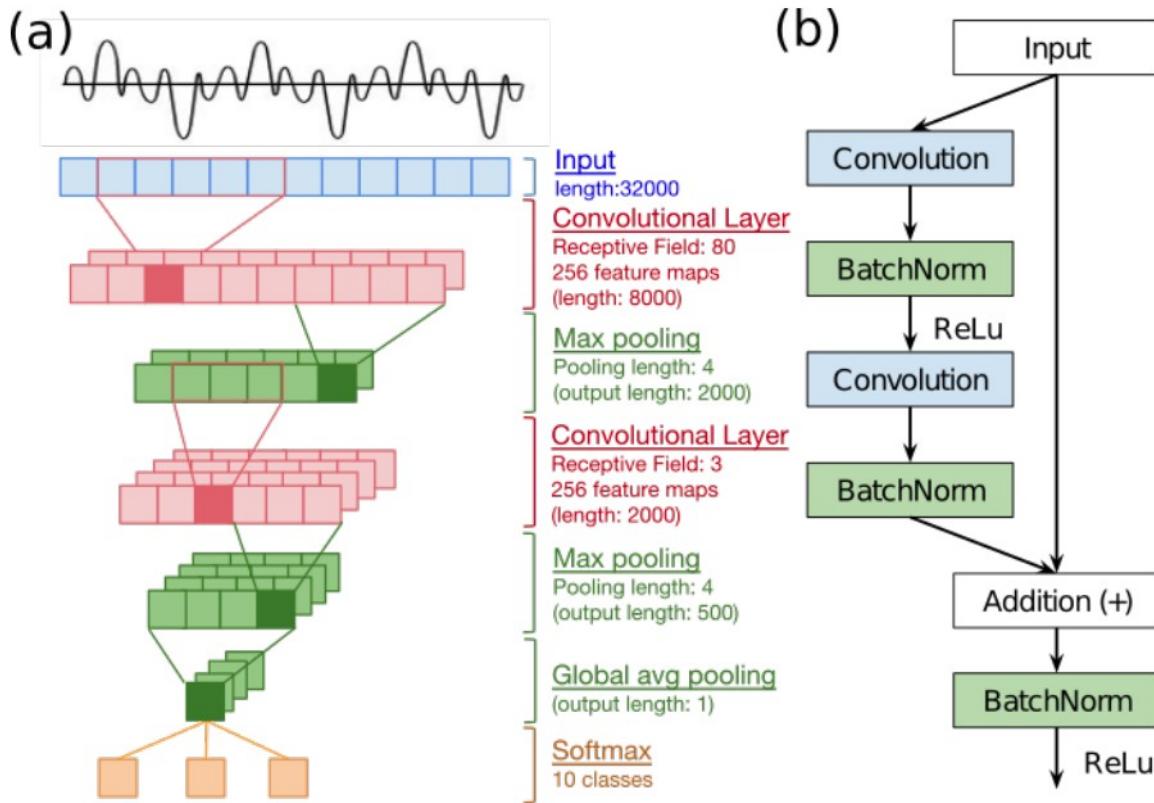
```
targets = torch.stack(targets)

return tensors, targets

batch_size = 256 #prima era 256 (accuratezza 74% con 3 epoch), 400
(71% con 3 epoch), 128(78% con 3 epoch)

#Always check which device you are using...
if device == "cuda":
    num_workers = 1
    pin_memory = True
else:
    num_workers = 0
    pin_memory = False

train_loader = torch.utils.data.DataLoader(
    train_set,
    batch_size=batch_size,
    shuffle=True,
    collate_fn=collate_fn,
    num_workers=num_workers,
    pin_memory=pin_memory,
)
test_loader = torch.utils.data.DataLoader(
    test_set,
    batch_size=batch_size,
    shuffle=False,
    drop_last=False,
    collate_fn=collate_fn,
    num_workers=num_workers,
    pin_memory=pin_memory,
)
```



We are ready to build our model. We are going to use an architecture from the literature. This is known as M5: <https://arxiv.org/pdf/1610.00087.pdf>

Report: study the different architectures inspect the difference between M3 and M5.

```

class M5(nn.Module):
    def __init__(self, n_input=1, n_output=35, stride=16,
n_channel=32):
        super().__init__()
        self.conv1 = nn.Conv1d(n_input, n_channel, kernel_size=80,
stride=stride)
        self.bn1 = nn.BatchNorm1d(n_channel)
        self.pool1 = nn.MaxPool1d(4)
        self.conv2 = nn.Conv1d(n_channel, n_channel, kernel_size=3)
        self.bn2 = nn.BatchNorm1d(n_channel)
        self.pool2 = nn.MaxPool1d(4)
        self.conv3 = nn.Conv1d(n_channel, 2 * n_channel,
kernel_size=3)
        self.bn3 = nn.BatchNorm1d(2 * n_channel)
        self.pool3 = nn.MaxPool1d(4)
        self.conv4 = nn.Conv1d(2 * n_channel, 2 * n_channel,
kernel_size=3)
        self.bn4 = nn.BatchNorm1d(2 * n_channel)
        self.pool4 = nn.MaxPool1d(4)
        self.fc1 = nn.Linear(2 * n_channel, n_output)
  
```

```

def forward(self, x):
    x = self.conv1(x)
    x = F.relu(self.bn1(x))
    x = self.pool1(x)
    x = self.conv2(x)
    x = F.relu(self.bn2(x))
    x = self.pool2(x)
    x = self.conv3(x)
    x = F.relu(self.bn3(x))
    x = self.pool3(x)
    x = self.conv4(x)
    x = F.relu(self.bn4(x))
    x = self.pool4(x)
    x = F.avg_pool1d(x, x.shape[-1])
    x = x.permute(0, 2, 1)
    x = self.fc1(x)
    return F.log_softmax(x, dim=2)

model = M5(n_input=transformed.shape[0], n_output=len(labels))
model.to(device)
print(model)

def count_parameters(model):
    return sum(p.numel() for p in model.parameters() if p.requires_grad)

n = count_parameters(model)
print("Number of parameters: %s" % n)

M5(
    (conv1): Conv1d(1, 32, kernel_size=(80,), stride=(16,))
    (bn1): BatchNorm1d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool1): MaxPool1d(kernel_size=4, stride=4, padding=0, dilation=1,
ceil_mode=False)
    (conv2): Conv1d(32, 32, kernel_size=(3,), stride=(1,))
    (bn2): BatchNorm1d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool2): MaxPool1d(kernel_size=4, stride=4, padding=0, dilation=1,
ceil_mode=False)
    (conv3): Conv1d(32, 64, kernel_size=(3,), stride=(1,))
    (bn3): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool3): MaxPool1d(kernel_size=4, stride=4, padding=0, dilation=1,
ceil_mode=False)
    (conv4): Conv1d(64, 64, kernel_size=(3,), stride=(1,)))

```

```

        (bn4): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (pool4): MaxPool1d(kernel_size=4, stride=4, padding=0, dilation=1,
ceil_mode=False)
        (fc1): Linear(in_features=64, out_features=35, bias=True)
)
Number of parameters: 26915

class M3(nn.Module):
    def __init__(self, n_input=1, n_output=35, stride=4,
n_channel=128):
        super().__init__()
        self.conv1 = nn.Conv1d(n_input, 2*n_channel, kernel_size=80,
stride=stride)
        self.bn1 = nn.BatchNorm1d(2*n_channel)
        self.pool1 = nn.MaxPool1d(4)
        self.conv2 = nn.Conv1d(2*n_channel, 2*n_channel,
kernel_size=3, stride=1, padding=1)
        self.bn2 = nn.BatchNorm1d(2*n_channel)
        self.pool2 = nn.MaxPool1d(4)
        self.fc1 = nn.Linear(2 * n_channel, n_output)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(self.bn1(x))
        x = self.pool1(x)
        x = self.conv2(x)
        x = F.relu(self.bn2(x))
        x = self.pool2(x)
        x = F.avg_pool1d(x, x.shape[-1])
        x = x.permute(0, 2, 1)
        x = self.fc1(x)
        return F.log_softmax(x, dim=2)

print('M3 model')
print('Model n_input', transformed.shape[0])
print('Model n_output', len(labels))
print('transformed :', transformed[0])
model_M3 = M3(n_input=transformed.shape[0], n_output=len(labels))
model_M3.to(device)
print(model_M3)

n = count_parameters(model_M3)
print("Number of parameters: %s" % n)

```

```

M3 model
Model n_input 1
Model n_output 35
transformed : tensor([-5.6026e-05, -2.9695e-04, -4.3415e-04, ..., -
3.9823e-05,
-6.2212e-05, 2.1935e-04])
M3(
  (conv1): Conv1d(1, 256, kernel_size=(80,), stride=(4,))
  (bn1): BatchNorm1d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool1): MaxPool1d(kernel_size=4, stride=4, padding=0, dilation=1,
ceil_mode=False)
  (conv2): Conv1d(256, 256, kernel_size=(3,), stride=(1,),
padding=(1,))
  (bn2): BatchNorm1d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool2): MaxPool1d(kernel_size=4, stride=4, padding=0, dilation=1,
ceil_mode=False)
  (fc1): Linear(in_features=256, out_features=35, bias=True)
)
Number of parameters: 227619

```

We are going to use an Adam optimizer and a StepLR reducing the learning rate after every 20 epochs of a factor of 10

```

optimizer = optim.Adam(model.parameters(), lr=0.01,
weight_decay=0.0001)
scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=20,
gamma=0.1) # reduce the learning after 20 epochs by a factor of 10

```

Let's define a "train" function for training our network.

```

def train(model, epoch, log_interval):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):

        data = data.to(device)
        target = target.to(device)

        # apply transform and model on whole batch directly on device
        data = transform(data)
        output = model(data)

        # negative log-likelihood for a tensor of size (batch x 1 x
n_output)
        loss = F.nll_loss(output.squeeze(), target)

        optimizer.zero_grad()
        loss.backward()

```

```

optimizer.step()

# print training stats
if batch_idx % log_interval == 0:
    print(f"Train Epoch: {epoch} [{batch_idx * len(data)}/{len(train_loader.dataset)} ({100. * batch_idx / len(train_loader):.0f}%)]\tLoss: {loss.item():.6f}")

# update progress bar
pbar.update(pbar_update)
# record loss
losses.append(loss.item())

```

We can build some functions to assess the performance of our network.

```

def number_of_correct(pred, target):
    # count number of correct predictions
    return pred.squeeze().eq(target).sum().item()

def get_likely_index(tensor):
    # find most likely label index for each element in the batch
    return tensor.argmax(dim=-1)

def test(model, epoch):
    model.eval()
    correct = 0
    for data, target in test_loader:

        data = data.to(device)
        target = target.to(device)

        # apply transform and model on whole batch directly on device
        data = transform(data)
        output = model(data)

        pred = get_likely_index(output)
        correct += number_of_correct(pred, target)

        # update progress bar
        pbar.update(pbar_update)

    print(f"\nTest Epoch: {epoch}\tAccuracy: {correct}/{len(test_loader.dataset)} ({100. * correct / len(test_loader.dataset):.0f}%)\n")

```

Let's train! The cell below trains our network for n_epoch. Also we use tqdm to build a nice completion bar (<https://tqdm.github.io/>)

```

#M3 train & validate
log_interval = 20
n_epoch=3

#Use Adam optimiser as used in paper. Weight decay set to 0.0001
#Initially train with learning rate of 0.01. Use a scheduler to
#decrease it to 0.001 during training after 20 epochs
optimiser = optim.Adam(model_M3.parameters(), lr=0.01,
weight_decay=0.0001)
scheduler = optim.lr_scheduler.StepLR(optimiser, step_size=20,
gamma=0.1) # reduce the learning after 20 epochs by a factor of 10

pbar_update = 1 / (len(train_loader) + len(test_loader))
losses = []

# The transform needs to live on the same device as the model and the
# data.
transform = transform.to(device)
with tqdm(total=n_epoch) as pbar:
    for epoch in range(1, n_epoch + 1):
        train(model_M3, epoch, log_interval)
        test(model_M3, epoch)
        scheduler.step()

# Let's plot the training loss versus the number of iteration.
# plt.plot(losses);
# plt.title("training loss");

0% | 0.002666666666666666/3 [00:02<53:15, 1066.20s/it]

Train Epoch: 1 [0/84843 (0%)] Loss: 3.615541

2%| | 0.05599999999999999/3 [00:23<20:31, 418.46s/it]

Train Epoch: 1 [5120/84843 (6%)] Loss: 3.604029

4%|| | 0.1093333333333328/3 [00:44<19:00, 394.67s/it]

Train Epoch: 1 [10240/84843 (12%)] Loss: 3.631300

5%|| | 0.1626666666666676/3 [01:03<17:02, 360.55s/it]

Train Epoch: 1 [15360/84843 (18%)] Loss: 3.585282

7%|| | 0.2160000000000033/3 [01:23<17:27, 376.13s/it]

Train Epoch: 1 [20480/84843 (24%)] Loss: 3.605152

9%|| | 0.2693333333333337/3 [01:44<20:06, 441.87s/it]

Train Epoch: 1 [25600/84843 (30%)] Loss: 3.615478

11%|| | 0.3226666666666667/3 [02:03<16:14, 364.03s/it]

```

```
Train Epoch: 1 [30720/84843 (36%)]    Loss: 3.611232
 13%|[■] | 0.3759999999999997/3 [02:23<18:24, 420.93s/it]
Train Epoch: 1 [35840/84843 (42%)]    Loss: 3.610683
 14%|[■] | 0.4293333333333273/3 [02:43<15:36, 364.21s/it]
Train Epoch: 1 [40960/84843 (48%)]    Loss: 3.601424
 16%|[■] | 0.4826666666666574/3 [03:03<16:28, 392.76s/it]
Train Epoch: 1 [46080/84843 (54%)]    Loss: 3.602304
 18%|[■] | 0.5359999999999995/3 [03:23<14:53, 362.62s/it]
Train Epoch: 1 [51200/84843 (60%)]    Loss: 3.611689
 20%|[■] | 0.5893333333333336/3 [03:43<14:57, 372.32s/it]
Train Epoch: 1 [56320/84843 (66%)]    Loss: 3.595884
 21%|[■] | 0.6426666666666677/3 [04:03<16:42, 425.35s/it]
Train Epoch: 1 [61440/84843 (72%)]    Loss: 3.610765
 23%|[■] | 0.69600000000000018/3 [04:24<14:55, 388.76s/it]
Train Epoch: 1 [66560/84843 (78%)]    Loss: 3.636536
 25%|[■] | 0.7493333333333336/3 [04:44<15:37, 416.42s/it]
Train Epoch: 1 [71680/84843 (84%)]    Loss: 3.581586
 27%|[■] | 0.8026666666666701/3 [05:05<13:48, 376.83s/it]
Train Epoch: 1 [76800/84843 (90%)]    Loss: 3.579956
 29%|[■] | 0.8560000000000042/3 [05:25<13:44, 384.71s/it]
Train Epoch: 1 [81920/84843 (96%)]    Loss: 3.618883
 33%|[■] | 1.0000000000000062/3 [06:15<11:10,
335.45s/it]/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:143: UserWarning: Detected call of `lr_scheduler.step()` before `optimizer.step()`. In PyTorch 1.1.0 and later, you should call them in the opposite order: `optimizer.step()` before `lr_scheduler.step()`. Failure to do this will result in PyTorch skipping the first value of the learning rate schedule. See more details at https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate
  warnings.warn("Detected call of `lr_scheduler.step()` before `optimizer.step()`. ")
```

Test Epoch: 1 Accuracy: 196/11005 (2%)

33% | [██████] | 1.002666666666728/3 [06:16<11:30, 345.93s/it]

Train Epoch: 2 [0/84843 (0%)] Loss: 3.564840

35% | [██████] | 1.0560000000000047/3 [06:36<14:06, 435.48s/it]

Train Epoch: 2 [5120/84843 (6%)] Loss: 3.600129

37% | [██████] | 1.1093333333333366/3 [06:56<11:39, 369.76s/it]

Train Epoch: 2 [10240/84843 (12%)] Loss: 3.580287

39% | [██████] | 1.1626666666666685/3 [07:16<12:50, 419.24s/it]

Train Epoch: 2 [15360/84843 (18%)] Loss: 3.581472

41% | [██████] | 1.2160000000000004/3 [07:36<10:44, 361.39s/it]

Train Epoch: 2 [20480/84843 (24%)] Loss: 3.602514

42% | [██████] | 1.2693333333333323/3 [07:56<11:45, 407.43s/it]

Train Epoch: 2 [25600/84843 (30%)] Loss: 3.594836

44% | [██████] | 1.3226666666666642/3 [08:17<10:32, 377.14s/it]

Train Epoch: 2 [30720/84843 (36%)] Loss: 3.600906

46% | [██████] | 1.3759999999999961/3 [08:37<10:04, 372.20s/it]

Train Epoch: 2 [35840/84843 (42%)] Loss: 3.617660

48% | [██████] | 1.429333333333328/3 [08:57<11:48, 450.78s/it]

Train Epoch: 2 [40960/84843 (48%)] Loss: 3.617874

49% | [██████] | 1.482666666666666/3 [09:17<09:16, 366.60s/it]

Train Epoch: 2 [46080/84843 (54%)] Loss: 3.600103

51% | [██████] | 1.5359999999999918/3 [09:37<10:18, 422.15s/it]

Train Epoch: 2 [51200/84843 (60%)] Loss: 3.611929

53% | [██████] | 1.5893333333333237/3 [09:56<08:33, 364.18s/it]

Train Epoch: 2 [56320/84843 (66%)] Loss: 3.617229

55% | [██████] | 1.6426666666666556/3 [10:17<08:42, 384.75s/it]

Train Epoch: 2 [61440/84843 (72%)] Loss: 3.624981

```
57%|███████ | 1.6959999999999875/3 [10:36<07:51, 361.61s/it]
Train Epoch: 2 [66560/84843 (78%)] Loss: 3.597152
58%|███████ | 1.749333333333194/3 [10:57<07:47, 373.53s/it]
Train Epoch: 2 [71680/84843 (84%)] Loss: 3.628718
60%|███████ | 1.8026666666666513/3 [11:16<08:19, 417.14s/it]
Train Epoch: 2 [76800/84843 (90%)] Loss: 3.600456
62%|███████ | 1.8559999999999832/3 [11:36<07:03, 370.35s/it]
Train Epoch: 2 [81920/84843 (96%)] Loss: 3.600364
65%|███████ | 1.9546666666666472/3 [12:12<05:58, 343.33s/it]

log_interval = 20
n_epoch = 3
#con 11 epoch e batch size=256 accuratezza 80%

pbar_update = 1 / (len(train_loader) + len(test_loader))
losses = []

# The transform needs to live on the same device as the model and the
# data.
transform = transform.to(device)
with tqdm(total=n_epoch) as pbar:
    for epoch in range(1, n_epoch + 1):
        train(model_M3, epoch, log_interval)
        test(model_M3, epoch)
        scheduler.step()

# Let's plot the training loss versus the number of iteration.
# plt.plot(losses);
# plt.title("training loss");

0%|          | 0.002666666666666666/3 [00:02<43:38, 873.45s/it]
Train Epoch: 1 [0/84843 (0%)] Loss: 3.621367
2%||         | 0.0559999999999999/3 [00:23<18:38, 379.83s/it]
Train Epoch: 1 [5120/84843 (6%)] Loss: 3.605422
4%||         | 0.1093333333333328/3 [00:43<21:34, 447.68s/it]
Train Epoch: 1 [10240/84843 (12%)] Loss: 3.640987
5%||         | 0.16266666666666676/3 [01:03<17:29, 370.02s/it]
Train Epoch: 1 [15360/84843 (18%)] Loss: 3.598700
```

7%|█ | 0.21600000000000033/3 [01:24<19:11, 413.56s/it]
Train Epoch: 1 [20480/84843 (24%)] Loss: 3.642989
9%|█ | 0.269333333333337/3 [01:43<16:41, 366.71s/it]
Train Epoch: 1 [25600/84843 (30%)] Loss: 3.604318
11%|█ | 0.3226666666666667/3 [02:04<17:07, 383.75s/it]
Train Epoch: 1 [30720/84843 (36%)] Loss: 3.598319
13%|█ | 0.3759999999999997/3 [02:24<19:07, 437.16s/it]
Train Epoch: 1 [35840/84843 (42%)] Loss: 3.603432
14%|█ | 0.4293333333333273/3 [02:44<15:56, 372.20s/it]
Train Epoch: 1 [40960/84843 (48%)] Loss: 3.593555
16%|█ | 0.4826666666666574/3 [03:05<17:47, 424.13s/it]
Train Epoch: 1 [46080/84843 (54%)] Loss: 3.606512
18%|█ | 0.5359999999999995/3 [03:25<15:25, 375.72s/it]
Train Epoch: 1 [51200/84843 (60%)] Loss: 3.625781
20%|█ | 0.5893333333333336/3 [03:46<15:54, 396.15s/it]
Train Epoch: 1 [56320/84843 (66%)] Loss: 3.630934
21%|█ | 0.6426666666666677/3 [04:07<17:33, 446.86s/it]
Train Epoch: 1 [61440/84843 (72%)] Loss: 3.594502
23%|█ | 0.6960000000000018/3 [04:27<14:21, 373.82s/it]
Train Epoch: 1 [66560/84843 (78%)] Loss: 3.607289
25%|█ | 0.7493333333333336/3 [04:48<15:45, 420.04s/it]
Train Epoch: 1 [71680/84843 (84%)] Loss: 3.594518
27%|█ | 0.8026666666666701/3 [05:08<13:33, 370.24s/it]
Train Epoch: 1 [76800/84843 (90%)] Loss: 3.603799
29%|█ | 0.8560000000000042/3 [05:29<13:45, 385.08s/it]
Train Epoch: 1 [81920/84843 (96%)] Loss: 3.599305
33%|█ | 1.0000000000000062/3 [06:19<11:21, 340.85s/it]

Test Epoch: 1 Accuracy: 241/11005 (2%)

33% | [██████] | 1.002666666666728/3 [06:20<11:40, 350.94s/it]

Train Epoch: 2 [0/84843 (0%)] Loss: 3.629742

35% | [██████] | 1.0560000000000047/3 [06:41<14:57, 461.82s/it]

Train Epoch: 2 [5120/84843 (6%)] Loss: 3.603289

37% | [██████] | 1.1093333333333366/3 [07:01<11:44, 372.40s/it]

Train Epoch: 2 [10240/84843 (12%)] Loss: 3.607699

39% | [██████] | 1.1626666666666685/3 [07:21<12:29, 408.18s/it]

Train Epoch: 2 [15360/84843 (18%)] Loss: 3.596784

41% | [██████] | 1.2160000000000004/3 [07:41<11:09, 375.05s/it]

Train Epoch: 2 [20480/84843 (24%)] Loss: 3.605073

42% | [██████] | 1.2693333333333323/3 [08:02<11:03, 383.23s/it]

Train Epoch: 2 [25600/84843 (30%)] Loss: 3.601493

44% | [██████] | 1.3226666666666642/3 [08:22<11:50, 423.47s/it]

Train Epoch: 2 [30720/84843 (36%)] Loss: 3.611003

46% | [██████] | 1.3759999999999961/3 [08:42<10:15, 378.99s/it]

Train Epoch: 2 [35840/84843 (42%)] Loss: 3.617035

48% | [██████] | 1.429333333333328/3 [09:03<10:59, 419.92s/it]

Train Epoch: 2 [40960/84843 (48%)] Loss: 3.618627

49% | [██████] | 1.482666666666666/3 [09:23<09:16, 367.01s/it]

Train Epoch: 2 [46080/84843 (54%)] Loss: 3.575739

51% | [██████] | 1.5359999999999918/3 [09:44<09:27, 387.90s/it]

Train Epoch: 2 [51200/84843 (60%)] Loss: 3.622574

53% | [██████] | 1.5893333333333237/3 [10:04<09:41, 412.54s/it]

Train Epoch: 2 [56320/84843 (66%)] Loss: 3.591389

55% | [██████] | 1.6426666666666556/3 [10:24<08:30, 376.35s/it]

Train Epoch: 2 [61440/84843 (72%)] Loss: 3.592136

57%|██████ | 1.6959999999999875/3 [10:45<09:19, 428.92s/it]
Train Epoch: 2 [66560/84843 (78%)] Loss: 3.612226
58%|██████ | 1.749333333333194/3 [11:04<07:41, 369.01s/it]
Train Epoch: 2 [71680/84843 (84%)] Loss: 3.633183
60%|██████ | 1.8026666666666513/3 [11:25<07:48, 391.63s/it]
Train Epoch: 2 [76800/84843 (90%)] Loss: 3.623835
62%|██████ | 1.8559999999999832/3 [11:44<06:56, 363.84s/it]
Train Epoch: 2 [81920/84843 (96%)] Loss: 3.618730
67%|██████ | 1.9999999999999793/3 [12:35<05:19, 319.15s/it]

Test Epoch: 2 Accuracy: 237/11005 (2%)

67%|██████ | 2.002666666666646/3 [12:36<05:38, 339.71s/it]
Train Epoch: 3 [0/84843 (0%)] Loss: 3.612593
69%|██████ | 2.0559999999999823/3 [12:56<05:55, 376.90s/it]
Train Epoch: 3 [5120/84843 (6%)] Loss: 3.631214
70%|██████ | 2.109333333333186/3 [13:16<06:36, 445.62s/it]
Train Epoch: 3 [10240/84843 (12%)] Loss: 3.556427
72%|██████ | 2.162666666666655/3 [13:36<05:04, 363.25s/it]
Train Epoch: 3 [15360/84843 (18%)] Loss: 3.622924
74%|██████ | 2.2159999999999913/3 [13:56<05:19, 408.11s/it]
Train Epoch: 3 [20480/84843 (24%)] Loss: 3.606650
76%|██████ | 2.269333333333277/3 [14:15<04:28, 367.58s/it]
Train Epoch: 3 [25600/84843 (30%)] Loss: 3.606331
77%|██████ | 2.322666666666664/3 [14:36<04:19, 383.05s/it]
Train Epoch: 3 [30720/84843 (36%)] Loss: 3.602437
79%|██████ | 2.3760000000000003/3 [14:55<04:00, 385.11s/it]
Train Epoch: 3 [35840/84843 (42%)] Loss: 3.591383
81%|██████ | 2.4293333333333367/3 [15:16<03:33, 373.59s/it]

```
Train Epoch: 3 [40960/84843 (48%)]    Loss: 3.618559
83%|[██████| 2.482666666666673/3 [15:36<03:50, 446.29s/it]
Train Epoch: 3 [46080/84843 (54%)]    Loss: 3.602975
85%|[██████| 2.53600000000000094/3 [15:55<02:51, 370.68s/it]
Train Epoch: 3 [51200/84843 (60%)]    Loss: 3.615752
86%|[██████| 2.589333333333457/3 [16:16<02:48, 409.56s/it]
Train Epoch: 3 [56320/84843 (66%)]    Loss: 3.596627
88%|[██████| 2.642666666666682/3 [16:35<02:08, 359.76s/it]
Train Epoch: 3 [61440/84843 (72%)]    Loss: 3.600069
90%|[██████| 2.6960000000000184/3 [16:55<02:00, 397.87s/it]
Train Epoch: 3 [66560/84843 (78%)]    Loss: 3.612092
92%|[██████| 2.749333333333547/3 [17:15<01:32, 367.99s/it]
Train Epoch: 3 [71680/84843 (84%)]    Loss: 3.638482
93%|[██████| 2.802666666666691/3 [17:35<01:13, 371.94s/it]
Train Epoch: 3 [76800/84843 (90%)]    Loss: 3.591575
95%|[██████| 2.8560000000000274/3 [17:55<01:04, 450.54s/it]
Train Epoch: 3 [81920/84843 (96%)]    Loss: 3.631505
100%|[██████| 2.997333333333687/3 [18:43<00:00,
322.51s/it]/usr/local/lib/python3.10/dist-packages/tqdm/std.py:636:
TqdmWarning: clamping frac to range [0, 1]
  full_bar = Bar(frac,
100%|[██████| 3.000000000000355/3 [18:45<00:00, 375.01s/it]
```

```
Test Epoch: 3    Accuracy: 241/11005 (2%)
```

```
def count_confusion_matrix(model, data_loader):
    model.eval()
    tp, fp, tn, fn = 0, 0, 0, 0
    with torch.no_grad():
        for data, target in data_loader:
            data = data.to(device)
            target = target.to(device)
            # apply transform and model on the whole batch directly on
```

```

the device
    data = transform(data)
    output = model(data)
    pred = get_likely_index(output)

    # Calculate TP, FP, TN, FN
    tp += ((pred == 1) & (target == 1)).sum().item()
    fp += ((pred == 1) & (target == 0)).sum().item()
    tn += ((pred == 0) & (target == 0)).sum().item()
    fn += ((pred == 0) & (target == 1)).sum().item()

return tp, fp, tn, fn

# Calculate counts for the training set
tp_train, fp_train, tn_train, fn_train =
count_confusion_matrix(model_M3, train_loader)

# Calculate counts for the test set
tp_test, fp_test, tn_test, fn_test = count_confusion_matrix(model_M3,
test_loader)

# Print counts
print("Training Set:")
print("True Positives (TP):", tp_train)
print("False Positives (FP):", fp_train)
print("True Negatives (TN):", tn_train)
print("False Negatives (FN):", fn_train)
print("\nTest Set:")
print("True Positives (TP):", tp_test)
print("False Positives (FP):", fp_test)
print("True Negatives (TN):", tn_test)
print("False Negatives (FN):", fn_test)

Training Set:
True Positives (TP): 0
False Positives (FP): 0
True Negatives (TN): 0
False Negatives (FN): 0

Test Set:
True Positives (TP): 0
False Positives (FP): 0
True Negatives (TN): 0
False Negatives (FN): 0

log_interval = 20
n_epoch = 3
#con 11 epoche e batch size=256 accuratezza 80%

pbar_update = 1 / (len(train_loader) + len(test_loader))

```

```
losses = []

# The transform needs to live on the same device as the model and the
# data.
transform = transform.to(device)
with tqdm(total=n_epoch) as pbar:
    for epoch in range(1, n_epoch + 1):
        train(model, epoch, log_interval)
        test(model, epoch)
        scheduler.step()

# Let's plot the training loss versus the number of iteration.
# plt.plot(losses);
# plt.title("training loss");

0%| 0.0026666666666666666666/3 [00:01<30:43, 615.04s/it]

Train Epoch: 1 [0/84843 (0%)] Loss: 3.790477

2%|| 0.05599999999999999/3 [00:19<16:21, 333.44s/it]

Train Epoch: 1 [5120/84843 (6%)] Loss: 3.048852

4%|| 0.1093333333333328/3 [00:38<18:45, 389.49s/it]

Train Epoch: 1 [10240/84843 (12%)] Loss: 2.570973

5%|| 0.1626666666666676/3 [00:56<15:51, 335.52s/it]

Train Epoch: 1 [15360/84843 (18%)] Loss: 2.277411

7%|| 0.2160000000000033/3 [01:14<14:59, 323.17s/it]

Train Epoch: 1 [20480/84843 (24%)] Loss: 2.050221

9%|| 0.2693333333333337/3 [01:33<15:56, 350.43s/it]

Train Epoch: 1 [25600/84843 (30%)] Loss: 1.863413

11%|| 0.3226666666666667/3 [01:52<14:30, 325.08s/it]

Train Epoch: 1 [30720/84843 (36%)] Loss: 1.911834

13%|| 0.3759999999999997/3 [02:11<15:27, 353.46s/it]

Train Epoch: 1 [35840/84843 (42%)] Loss: 1.695160

14%|| 0.4293333333333273/3 [02:29<14:15, 332.72s/it]

Train Epoch: 1 [40960/84843 (48%)] Loss: 1.840236

16%|| 0.48266666666666574/3 [02:48<15:30, 369.52s/it]

Train Epoch: 1 [46080/84843 (54%)] Loss: 1.641245
```

18% | [■] | 0.5359999999999995/3 [03:06<13:25, 326.80s/it]
Train Epoch: 1 [51200/84843 (60%)] Loss: 1.559800
20% | [■] | 0.5893333333333336/3 [03:25<16:59, 422.75s/it]
Train Epoch: 1 [56320/84843 (66%)] Loss: 1.461332
21% | [■] | 0.6426666666666677/3 [03:43<13:08, 334.64s/it]
Train Epoch: 1 [61440/84843 (72%)] Loss: 1.356640
23% | [■] | 0.69600000000000018/3 [04:01<14:05, 367.02s/it]
Train Epoch: 1 [66560/84843 (78%)] Loss: 1.059988
25% | [■] | 0.7493333333333336/3 [04:20<12:46, 340.72s/it]
Train Epoch: 1 [71680/84843 (84%)] Loss: 1.160042
27% | [■] | 0.8026666666666701/3 [04:38<12:01, 328.17s/it]
Train Epoch: 1 [76800/84843 (90%)] Loss: 1.009504
29% | [■] | 0.8560000000000042/3 [04:57<12:37, 353.09s/it]
Train Epoch: 1 [81920/84843 (96%)] Loss: 1.138900
33% | [■] | 1.0000000000000062/3 [05:44<13:03, 391.92s/it]

Test Epoch: 1 Accuracy: 7137/11005 (65%)

33% | [■] | 1.0026666666666728/3 [05:45<12:41, 381.26s/it]
Train Epoch: 2 [0/84843 (0%)] Loss: 1.034717
35% | [■] | 1.0560000000000047/3 [06:03<10:47, 333.19s/it]
Train Epoch: 2 [5120/84843 (6%)] Loss: 1.092970
37% | [■] | 1.1093333333333366/3 [06:22<12:14, 388.70s/it]
Train Epoch: 2 [10240/84843 (12%)] Loss: 1.062043
39% | [■] | 1.1626666666666685/3 [06:40<10:13, 333.71s/it]
Train Epoch: 2 [15360/84843 (18%)] Loss: 1.036001
41% | [■] | 1.2160000000000004/3 [06:58<10:00, 336.41s/it]
Train Epoch: 2 [20480/84843 (24%)] Loss: 0.923860
42% | [■] | 1.2693333333333323/3 [07:17<09:42, 336.69s/it]

```
Train Epoch: 2 [25600/84843 (30%)] Loss: 1.022502
 44%|██████ | 1.3226666666666642/3 [07:35<09:05, 325.39s/it]
Train Epoch: 2 [30720/84843 (36%)] Loss: 1.059344
 46%|██████ | 1.3759999999999961/3 [07:54<09:34, 354.01s/it]
Train Epoch: 2 [35840/84843 (42%)] Loss: 1.076034
 48%|██████ | 1.42933333333328/3 [08:12<08:32, 326.27s/it]
Train Epoch: 2 [40960/84843 (48%)] Loss: 1.231857
 49%|██████ | 1.482666666666666/3 [08:31<09:42, 384.10s/it]
Train Epoch: 2 [46080/84843 (54%)] Loss: 0.969934
 51%|██████ | 1.5359999999999918/3 [08:49<08:03, 329.96s/it]
Train Epoch: 2 [51200/84843 (60%)] Loss: 0.991519
 53%|██████ | 1.589333333333237/3 [09:07<09:45, 415.20s/it]
Train Epoch: 2 [56320/84843 (66%)] Loss: 0.976233
 55%|██████ | 1.6426666666666556/3 [09:25<07:39, 338.86s/it]
Train Epoch: 2 [61440/84843 (72%)] Loss: 0.748306
 57%|██████ | 1.6959999999999875/3 [09:44<07:42, 354.40s/it]
Train Epoch: 2 [66560/84843 (78%)] Loss: 0.796559
 58%|██████ | 1.749333333333194/3 [10:02<07:12, 346.00s/it]
Train Epoch: 2 [71680/84843 (84%)] Loss: 0.826512
 60%|██████ | 1.8026666666666513/3 [10:20<06:28, 324.31s/it]
Train Epoch: 2 [76800/84843 (90%)] Loss: 0.911608
 62%|██████ | 1.8559999999999832/3 [10:39<06:40, 349.74s/it]
Train Epoch: 2 [81920/84843 (96%)] Loss: 0.790906
 67%|██████ | 1.9999999999999793/3 [11:26<06:16, 376.06s/it]

Test Epoch: 2 Accuracy: 7722/11005 (70%)

 67%|██████ | 2.002666666666646/3 [11:27<06:26, 387.56s/it]
Train Epoch: 3 [0/84843 (0%)] Loss: 0.908686
```

69% | ██████████ | 2.0559999999999823/3 [11:45<05:13, 331.74s/it]
Train Epoch: 3 [5120/84843 (6%)] Loss: 0.817632
70% | ██████████ | 2.109333333333186/3 [12:03<05:38, 380.44s/it]
Train Epoch: 3 [10240/84843 (12%)] Loss: 0.907780
72% | ██████████ | 2.162666666666655/3 [12:22<04:43, 338.60s/it]
Train Epoch: 3 [15360/84843 (18%)] Loss: 0.862680
74% | ██████████ | 2.2159999999999913/3 [12:40<04:14, 324.29s/it]
Train Epoch: 3 [20480/84843 (24%)] Loss: 0.856575
76% | ██████████ | 2.269333333333277/3 [12:59<04:15, 349.88s/it]
Train Epoch: 3 [25600/84843 (30%)] Loss: 0.883612
77% | ██████████ | 2.3226666666666664/3 [13:17<03:40, 324.99s/it]
Train Epoch: 3 [30720/84843 (36%)] Loss: 0.708188
79% | ██████████ | 2.3760000000000003/3 [13:36<03:48, 365.71s/it]
Train Epoch: 3 [35840/84843 (42%)] Loss: 0.759057
81% | ██████████ | 2.4293333333333367/3 [13:53<03:06, 325.98s/it]
Train Epoch: 3 [40960/84843 (48%)] Loss: 0.809799
83% | ██████████ | 2.482666666666673/3 [14:12<03:26, 398.41s/it]
Train Epoch: 3 [46080/84843 (54%)] Loss: 0.994512
85% | ██████████ | 2.5360000000000094/3 [14:30<02:33, 330.75s/it]
Train Epoch: 3 [51200/84843 (60%)] Loss: 0.572577
86% | ██████████ | 2.5893333333333457/3 [14:49<02:33, 372.84s/it]
Train Epoch: 3 [56320/84843 (66%)] Loss: 0.905681
88% | ██████████ | 2.642666666666682/3 [15:07<02:01, 340.56s/it]
Train Epoch: 3 [61440/84843 (72%)] Loss: 0.923929
90% | ██████████ | 2.6960000000000184/3 [15:25<01:39, 325.70s/it]
Train Epoch: 3 [66560/84843 (78%)] Loss: 0.722016
92% | ██████████ | 2.7493333333333547/3 [15:44<01:26, 343.49s/it]
Train Epoch: 3 [71680/84843 (84%)] Loss: 0.797123

```

93%|██████████| 2.802666666666691/3 [16:02<01:06, 335.05s/it]
Train Epoch: 3 [76800/84843 (90%)] Loss: 0.744030
95%|██████████| 2.856000000000000274/3 [16:21<00:52, 362.85s/it]
Train Epoch: 3 [81920/84843 (96%)] Loss: 0.868759
100%|██████████| 3.0000000000000355/3 [17:07<00:00, 342.54s/it]

Test Epoch: 3 Accuracy: 8197/11005 (74%)

```

```

def count_confusion_matrix(model, data_loader):
    model.eval()
    tp, fp, tn, fn = 0, 0, 0, 0
    with torch.no_grad():
        for data, target in data_loader:
            data = data.to(device)
            target = target.to(device)
            # apply transform and model on the whole batch directly on
the device
            data = transform(data)
            output = model(data)
            pred = get_likely_index(output)

            # Calculate TP, FP, TN, FN
            tp += ((pred == 1) & (target == 1)).sum().item()
            fp += ((pred == 1) & (target == 0)).sum().item()
            tn += ((pred == 0) & (target == 0)).sum().item()
            fn += ((pred == 0) & (target == 1)).sum().item()

    return tp, fp, tn, fn

# Calculate counts for the training set
tp_train, fp_train, tn_train, fn_train = count_confusion_matrix(model,
train_loader)

# Calculate counts for the test set
tp_test, fp_test, tn_test, fn_test = count_confusion_matrix(model,
test_loader)

# Print counts
print("Training Set:")
print("True Positives (TP):", tp_train)
print("False Positives (FP):", fp_train)
print("True Negatives (TN):", tn_train)
print("False Negatives (FN):", fn_train)
print("\nTest Set:")

```

```
print("True Positives (TP):", tp_test)
print("False Positives (FP):", fp_test)
print("True Negatives (TN):", tn_test)
print("False Negatives (FN):", fn_test)
```

Training Set:

```
True Positives (TP): 8173
False Positives (FP): 6039
True Negatives (TN): 6496
False Negatives (FN): 6259
```

Test Set:

```
True Positives (TP): 21191
False Positives (FP): 1292
True Negatives (TN): 11272
False Negatives (FN): 1404
```

Report: inspect performance on the different classes. Refine the model and the training strategy to increase them.

Finally, we can use our model for making commands' classification. See the following function.

```
def predict(tensor):
    # Use the model to predict the label of the waveform
    tensor = tensor.to(device)
    tensor = transform(tensor)
    tensor = model(tensor.unsqueeze(0))
    tensor = get_likely_index(tensor)
    tensor = index_to_label(tensor.squeeze())
    return tensor

waveform, sample_rate, utterance, *_ = train_set[-1]
play_audio(waveform, sample_rate)

print(f"Expected: {utterance}. Predicted: {predict(waveform)}.")

<IPython.lib.display.Audio object>

Expected: zero. Predicted: zero.
```

Report Suggestions:

1. Understand M5 architecture, compare with M3 (see reference paper).
2. Improve model performance (74% accuracy...can we boost it? How?)
3. Assess the performance on the single classes (confusion matrix). What are the main findings?
4. How can we refine this model?